





ACTION CHUNKING

Delayed policies

Predict the current action a_t based on a past state s_{t-i}



Better capture human non-Markovianity

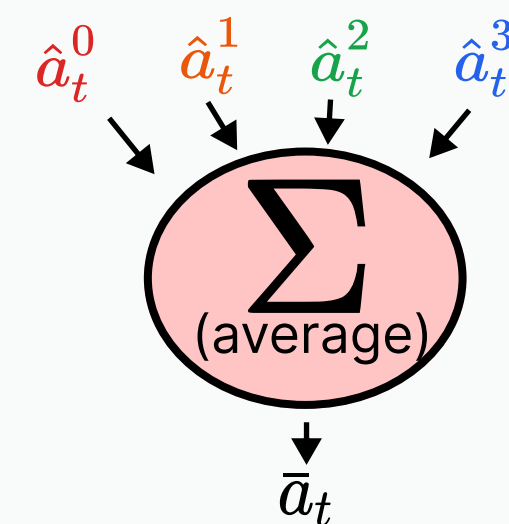
Less compounding errors

+

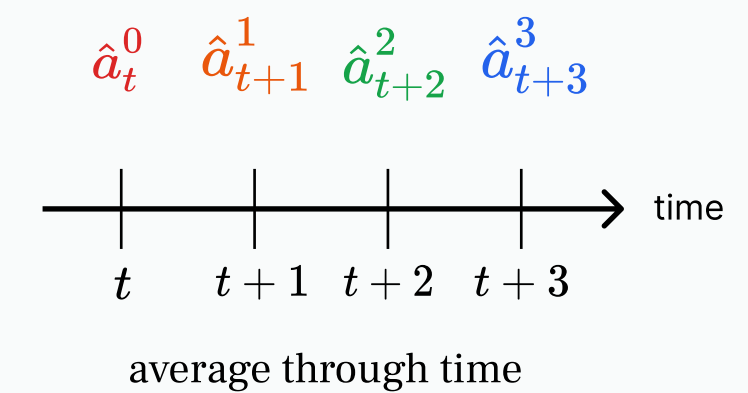
Ensembling

Aggregate action predictions from multiple delayed policies

Explicit (eg, ACT)



Implicit (eg, Diffusion Policy)



Lower generalization error

More stable predictions